

Assignment 1

1. (5) Consider the matrix $H(v) = \hat{1} - 2vv^T$, where v is a unit column vector. What is the rank of the matrix $H(v)$? Prove that it is orthogonal.
2. (10) Prove the following inequalities and provide examples of x and A when they turn into equalities:
 - $\|x\|_2 \leq \sqrt{m} \|x\|_\infty$
 - $\|A\|_\infty \leq \sqrt{n} \|A\|_2$

where x is a vector of m components and A is $m \times n$ matrix.

3. (5) Assuming u and v are m -vectors, consider the matrix $A = 1 + uv^T$ which is a rank-one perturbation of identity. Can it be singular? Assuming it is not, compute its inverse. You may look for it in a form of $A^{-1} = 1 + \alpha uv^T$ for some scalar α and evaluate α .
4. (5) Prove that for any unitary matrix U one has $\|UA\|_F = \|AU\|_F = \|A\|_F$.
5. (10) In this exercise your goal will be to study and speed up an implementation of [K-means algorithm](#). In the notebook `kmeans.ipynb`, you can find a naive implementation. Explore the code, make sure you understand it. You will find there two functions `dist_i` and `dist_ij` which are (on purpose) implemented in a rather inefficient way. Improve them by getting rid of the loops in the favor of a proper numpy vectorized implementation and measure the speed-up of the full algorithm for $N = 10000$.
6. (10) Some things just can not be vectorized but still can be sped up compared to naive implementation. For example, consider computation of the [Hofstadter-Conway sequence](#) $a(n)$ such that $a(1) = 1$, $a(2) = 1$ and

$$a(n) = a(a(n-1)) + a(n - a(n-1)), \quad n > 2 \quad (1)$$

Write three functions, computing the sequence up to n -th element in three ways: i) pre-allocating `numpy` array and filling it using `for` loop, ii) cumulatively appending `python` list and converting it to `numpy` array, iii) same as i) but compiled (`jit`) version. Time the resulting implementations and conclude which is preferable. With the optimal one, compute $a(10^8)$.

7. (15*) Consider a function mapping six tensors to one tensor: $Z(\lambda^{(1)}, \lambda^{(2)}, \lambda^{(3)}, \Gamma^{(1)}, \Gamma^{(2)}, U)$, with

$$Z_{ahij} = \sum_{bcdefg} \lambda^{(1)}_{ab} \Gamma^{(1)}_{cbd} \lambda^{(2)}_{de} \Gamma^{(2)}_{feg} \lambda^{(3)}_{gh} U_{ijcf}. \quad (2)$$

Assume that all indices of the tensors appearing above take values from 1 to χ . Running the numerical experiments, explore the values of χ in the range 3–50 (from slowest to fastest implementation).

- In the notebook `convolution.ipynb` you may find implemented a *stupid* way to compute this convolution, which takes $\chi^4 \times \chi^6 = \chi^{10}$ flops. In fact, this can be computed much faster!
- Using the function `numpy.einsum` (its crucial to use the `optimize` argument), you can actually achieve a much faster implementation. In order to understand what it is doing under the hood, explore the function `numpy.einsum_path`. What is the minimal number of flops required for computation of Z ?
- Using the understanding of the output of `numpy.einsum_path`, implement an algorithm to compute Z , which is as effective as `numpy.einsum`, but relying only on more elementary `numpy.dot` and `numpy.tensor_dot`.